

General Field Presentation

Cliff Notes

Sahara Kaul

Kelsey Pattison

Ahmed Sajid

CMPUT 414, Winter 2026

Purpose

Quick live-delivery notes. These are for glancing at during the talk, not reading word-for-word.

Speaker split

Kelsey: introduction and core challenges

Sarah: early generation, VAEs, RAVE, MoVE

Ahmed: modern controllable synthesis, diffusion, field trends, conclusion

Kelsey

Style transfer: what is it?

- Start with image analogy: preserve scene structure, change rendering style.
- In music: preserve content, change timbre, instrumentation, texture, articulation, production character.
- Mention demo if using one.

Transition: define content versus style clearly.

Style vs content

- Content: melody, rhythm, notes, harmonic path, phrase structure.
- Style: timbre, instrumentation, texture, dynamics, expressive details.
- Same content can sound completely different under a new style.

Transition: now explain why this is harder in music than it sounds.

Why music style transfer is hard

- Music is multilevel: melody, harmony, texture, production, room, long-range structure all mixed together.
- Weak systems create a “coat of paint” result instead of true re-authoring.
- Genre is broader than timbre: conventions, groove, production norms, listener expectations.

Transition: one of the key recurring ideas is representation.

Representations

- Raw waveforms, spectrograms / log-mel, MIDI, learned latent spaces.
- Each representation exposes some structure and hides other information.
- Representation choice matters as much as architecture choice.

Transition: briefly explain STFT and MIDI as anchor examples.

STFT and MIDI

- Fourier Transform: what frequencies exist overall.
- STFT: what frequencies exist over time, gives a spectrogram.
- MIDI is easy to manipulate but weak for final timbral realism.

Transition: this leads to the central field question.

Central field question

- How do we change the style or genre of a song while keeping the musical identity recognizable?

Transition: answer that by naming the field bottlenecks.

Core challenges

- Representation problem.
- Disentanglement problem.
- Realism problem.
- Long-form coherence problem.
- These bottlenecks give a better field structure than just listing models.

Transition: Sarah picks up with the early audio-generation story.

Sarah

WaveNet

- Early milestone in realistic neural audio.
- Autoregressive and slow: one sample at a time.
- Important historically, but not a practical direct answer to style transfer.

Transition: this motivates GANs and parallel generation.

Background about GANs

- Generator makes fake samples; discriminator spots fake vs real.
- Often sharper than plain autoencoder outputs.
- Training is unstable and highly dependent on representation choice.

Transition: first major audio GAN example is WaveGAN.

WaveGAN

- Parallel raw-audio generation, 1D convolutions, phase shuffle.
- Limits: phase precession and short fixed-length output.
- Shows raw waveform GANs are possible, but not enough for robust style transfer.

Transition: GANSynth improves the working space.

GANSynth

- Generate log-magnitude + instantaneous-frequency spectrograms instead of raw waveform.
- More stable than raw phase.
- Main lesson: representation choice matters.

Transition: summarize why early generation is only background.

Why early generation is mainly background

- WaveNet, WaveGAN, and GANSynth explain how neural audio generation developed.
- But they do not solve content preservation while changing style.
- Combined takeaway: generation is possible, representation matters, true transfer remains unsolved.

Transition: now move into latent structured audio.

VAEs

- Useful because they expose a latent space.
- Problems: posterior collapse, weak bottleneck, sequential complexity.
- Latent-space idea still matters for style transfer.

Transition: first talk about RAVE as practical latent audio, then MoVE as more transfer-specific.

RAVE

- Not a strong disentanglement model, but makes high-quality latent audio practical.
- Uses spectral losses, latent sequences, and adversarial sharpening.
- Best framed as latent audio infrastructure: encode, manipulate, resynthesize.

Transition: now move to a model that is more directly about transfer.

MoVE

- Explicitly about many-to-many musical timbre transfer.
- Uses VAE latent space, FiLM conditioning, MMD objective.
- More transfer-specific than RAVE, but still mainly timbre transfer rather than full genre remastering.

Transition: this latent section matters because it gets much closer to actual style transfer than the older GAN section.

Ahmed

Modern controllable synthesis

- New question: once we have a usable representation and some notion of style, how do we render convincing controllable audio?

Transition: start with BigVGAN and FiLM because they do different jobs.

BigVGAN

- Universal neural vocoder.
- Converts intermediate representation to waveform audio.
- Role in the story: realism stage, not transfer logic.

Transition: FiLM is about conditioning, not waveform realism.

FiLM

- Feature-wise Linear Modulation.
- Style vector scales and shifts hidden features.
- Good way to inject style information throughout the network.

Transition: now move into diffusion itself.

DDPM

- Forward noising process, learned reverse denoising process.
- More stable than many adversarial setups.
- Important because gradual refinement gives more control.

Transition: classifier-free guidance and SDEdit make that control more explicit.

Classifier-Free Guidance

- Guidance scale w controls how strongly generation follows the target condition.
- Lower guidance preserves naturalness; higher guidance pushes harder toward style.

Transition: SDEdit is the editing version of the same idea.

SDEdit

- Start from a noised source instead of pure noise.
- Lets the model preserve the source skeleton while still changing the sound world.

Transition: this sets up AudioLDM.

AudioLDM

- Latent diffusion for audio.
- Important because diffusion happens in a learned audio latent space instead of raw waveform space.
- Strong modern direction for controllable audio transformation and re-authoring.

Transition: now summarize why diffusion is the modern direction.

Why diffusion is the modern generative direction

- Progressive generation.
- Strong conditioning.
- Explicit guidance control.
- Source anchoring through SDEdit-style editing.

Transition: end by summarizing the overall field movement.

Overview of field trends

- More useful arc: disentanglement, controllable latent audio, diffusion-based re-authoring, long-form coherence.
- Field moved from raw generation toward structured representations and explicit control.

Transition: final lesson.

Conclusion

- Music style transfer is not just an audio-generation problem.
- It is also a structure-preservation, representation, conditioning, synthesis, and coherence problem.
- Clean ending: the field is moving from surface transfer toward structure-aware genre reconstruction.